

KASPERSKY



NETCAP: A framework for secure and concurrent network traffic analysis

\$ whoami

Student from the LMU Munich

Interests:

- Network Security Monitoring & Anomaly Detection
- Machine Learning
- Programming (Golang, C / C++ / ObjC, Swift, Haskell, Python, Rust)
- Hardware & Software Security
- Reverse Engineering
- Penetration Testing



Why network security monitoring?

Network environments have many vulnerable components

Hackers will always find a way in - “low hanging fruit”

Be prepared for the worst case!

Why anomaly detection?

Number of signatures is exploding (Kaspersky: 330 000 new samples per day)

New threats cannot be detected by signatures

Existing malware can be obfuscated to be fully undetectable again

Problems?

“The monitor will be attacked.”

–Vern Paxson, 1998

A major problem for data collection

memory safety

All existing frameworks are written in C or C++

Parsing network protocols is complex.

Mistakes happen a lot.

MITRE CVE results for Snort IDS

CVE-2016-1463	Cisco FireSIGHT System Software 5.3.0, 5.3.1, 5.4.0, 6.0, and 6.0.1 allows remote attackers to bypass Snort rules via crafted parameters in the header of an HTTP packet, aka Bug ID CSCuz20737.
CVE-2016-1417	Untrusted search path vulnerability in Snort 2.9.7.0-WIN32 allows remote attackers to execute arbitrary code and conduct DLL hijacking attacks via a Trojan horse tcapi.dll that is located in the same folder on a remote file share as a pcap file that is being processed.
CVE-2015-6427	Cisco FireSIGHT Management Center allows remote attackers to bypass the HTTP attack detection feature and avoid triggering Snort IDS rules via an SSL session that is mishandled after decryption, aka Bug ID CSCux53437.
CVE-2014-4695	Multiple open redirect vulnerabilities in the Snort package before 3.0.13 for pfSense through 2.1.4 allow remote attackers to redirect users to arbitrary web sites and conduct phishing attacks via (1) the referer parameter to snort_rules_flowbits.php or (2) the returl parameter to snort_select_alias.php.
CVE-2014-4693	Multiple cross-site scripting (XSS) vulnerabilities in the Snort package before 3.0.13 for pfSense through 2.1.4 allow remote attackers to inject arbitrary web script or HTML via (1) the eng parameter to snort_import_aliases.php or (2) unspecified variables to snort_select_alias.php.
CVE-2009-4211	The U.S. Defense Information Systems Agency (DISA) Security Readiness Review (SRR) script for the Solaris x86 platform executes files in arbitrary directories as root for filenames equal to (1) java, (2) openssl, (3) php, (4) snort, (5) tshark, (6) vncserver, or (7) wireshark, which allows local users to gain privileges via a Trojan horse program.
CVE-2009-3641	Snort before 2.8.5.1, when the -v option is enabled, allows remote attackers to cause a denial of service (application crash) via a crafted IPv6 packet that uses the (1) TCP or (2) ICMP protocol.
CVE-2007-0251	Integer underflow in the DecodeGRE function in src/decode.c in Snort 2.6.1.2 allows remote attackers to trigger dereferencing of certain memory locations via crafted GRE packets, which may cause corruption of log files or writing of sensitive information into log files.
CVE-2006-6931	Algorithmic complexity vulnerability in Snort before 2.6.1, during predicate evaluation in rule matching for certain rules, allows remote attackers to cause a denial of service (CPU consumption and detection outage) via crafted network traffic, aka a "backtracking attack."
CVE-2006-5276	Stack-based buffer overflow in the DCE/RPC preprocessor in Snort before 2.6.1.3, and 2.7 before beta 2; and Sourcefire Intrusion Sensor; allows remote attackers to execute arbitrary code via crafted SMB traffic.
CVE-2006-2769	The HTTP Inspect preprocessor (http_inspect) in Snort 2.4.0 through 2.4.4 allows remote attackers to bypass "uricontent" rules via a carriage return (\r) after the URL and before the HTTP declaration.

MITRE CVE results for Bro IDS

Search Results

There are **6** CVE entries that match your search.

Name	Description
CVE-2018-17019	In Bro through 2.5.5, there is a DoS in IRC protocol names command parsing in analyzer/protocol/irc/IRC.cc.
CVE-2018-16807	In Bro through 2.5.5, there is a memory leak potentially leading to DoS in scripts/base/protocols/krb/main.bro in the Kerberos protocol parser.
CVE-2017-1000458	Bro before Bro v2.5.2 is vulnerable to an out of bounds write in the ContentLine analyzer allowing remote attackers to cause a denial of service (crash) and possibly other exploitation.
CVE-2015-1522	analyzer/protocol/dnp3/DNP3.cc in Bro before 2.3.2 does not reject certain non-zero values of a packet length, which allows remote attackers to cause a denial of service (buffer overflow or buffer over-read) via a crafted DNP3 packet.
CVE-2015-1521	analyzer/protocol/dnp3/DNP3.cc in Bro before 2.3.2 does not properly handle zero values of a packet length, which allows remote attackers to cause a denial of service (buffer overflow or buffer over-read if NDEBUG; otherwise assertion failure) via a crafted DNP3 packet.
CVE-2007-0186	Multiple cross-site scripting (XSS) vulnerabilities in F5 FirePass SSL VPN allow remote attackers to inject arbitrary web script or HTML via (1) the xcho parameter to my.logon.php3; the (2) topblue, (3) midblue, (4) wtopblue, and certain other Custom color parameters in a per action to vdesk/admincon/index.php; the (5) h321, (6) h311, (7) h312, and certain other Front Door custom text color parameters in a per action to vdesk/admincon/index.php; the (8) ua parameter in a bro action to vdesk/admincon/index.php; the (9) app_param and (10) app_name parameters to webyfiers.php; (11) double eval functions; (12) JavaScript contained in an <FP_DO_NOT_TOUCH> element; and (13) the vhost parameter to my.activation.php. NOTE: it is possible that this candidate overlaps CVE-2006-3550.

Bro 2.5.5 primarily addresses security issues.

- Fix array bounds checking in BinPAC: for arrays that are fields within a record, the bounds check was based on a pointer to the start of the record rather than the start of the array field, potentially resulting in a buffer over-read.
- Fix SMTP command string comparisons: the number of bytes compared was based on the user-supplied string length and can lead to incorrect matches. e.g. giving a command of "X" incorrectly matched "X-ANONYMOUSTLS" (and an empty commands match anything).

The following changes address potential vectors for Denial of Service reported by Christian Titze & Jan Grashöfer of Karlsruhe Institute of Technology:

- "Weird" events are now generally suppressed/sampled by default according to some tunable parameters:
 - Weird::sampling_whitelist
 - Weird::sampling_threshold
 - Weird::sampling_rate
 - Weird::sampling_duration

Those options can be changed if one needs the previous behavior of a "net_weird", "flow_weird", or "conn_weird" event being raised for every single event. Otherwise, there is a new weird_stats.log which contains concise summaries of weird counts per type per time period and the original weird.log may not differ much either, except in the cases where a particular weird type exceeds the sampling threshold. These changes help improve performance issues resulting from excessive numbers of weird events.

- Improved handling of empty lines in several text protocol analyzers that can cause performance issues when seen in long sequences.
- Add 'smtp_excessive_pending_cmds' weird which serves as a notification for when the "pending

!Problem?



To the rescue!

NETCAP

Traffic Analysis Framework

How does it work?

Uses gopacket library for decoding packets in pure Go

Generates audit records as compressed protocol buffers

Concurrent design: worker pool, each audit record written to a separate file

-> enables consuming applications to implement a concurrent design as well

Why Protocol Buffers?

Type safe structured data - can represent complex nested structures

Goal: Data Accessibility. Generated Data is available in almost any programming language

Huge landscape of frameworks for machine learning, in many interesting languages
(R, Scala, Haskell...)

see: <https://github.com/josephmisiti/awesome-machine-learning>

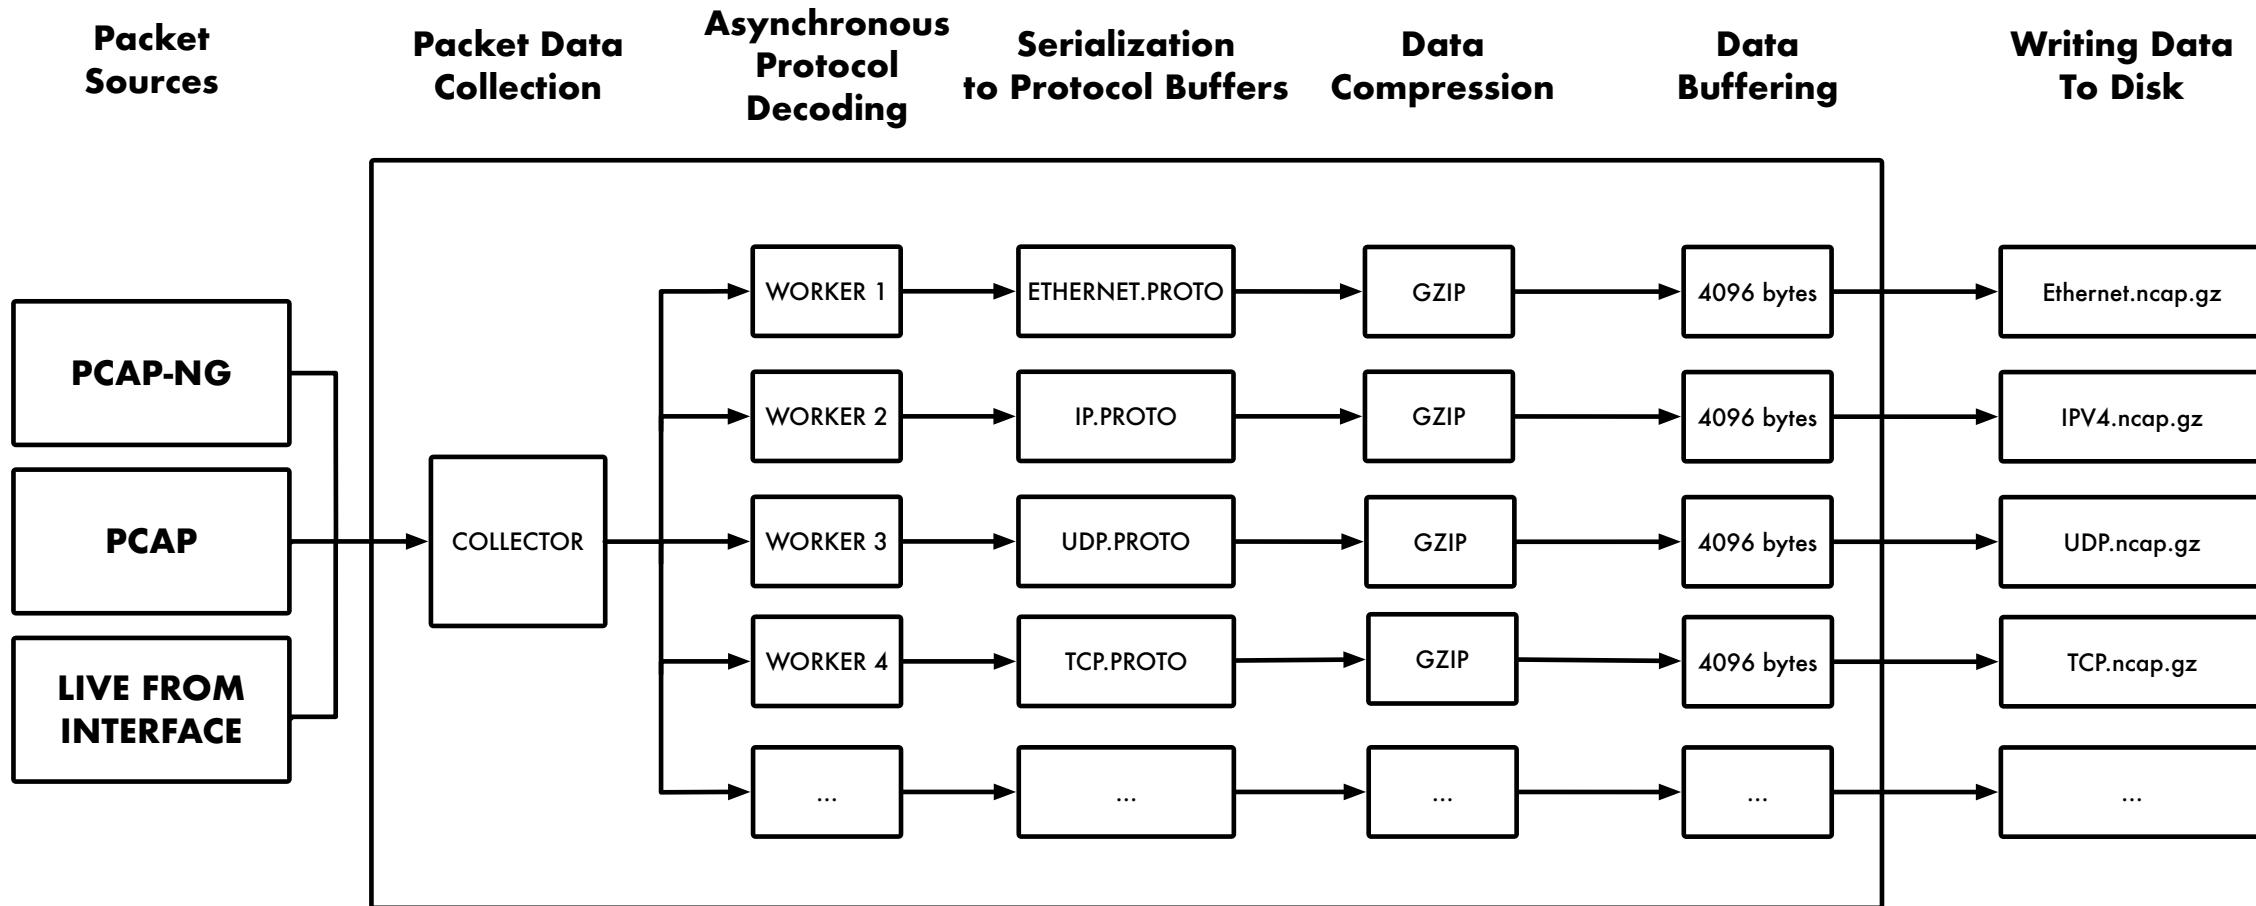
What audit records are generated?

- + TLS (Client Hello Msg + Ja3)
- + LinkFlow
- + NetworkFlow
- + TransportFlow
- + HTTP
- + Flow (unidirectional)
- + Connection (bidirectional)

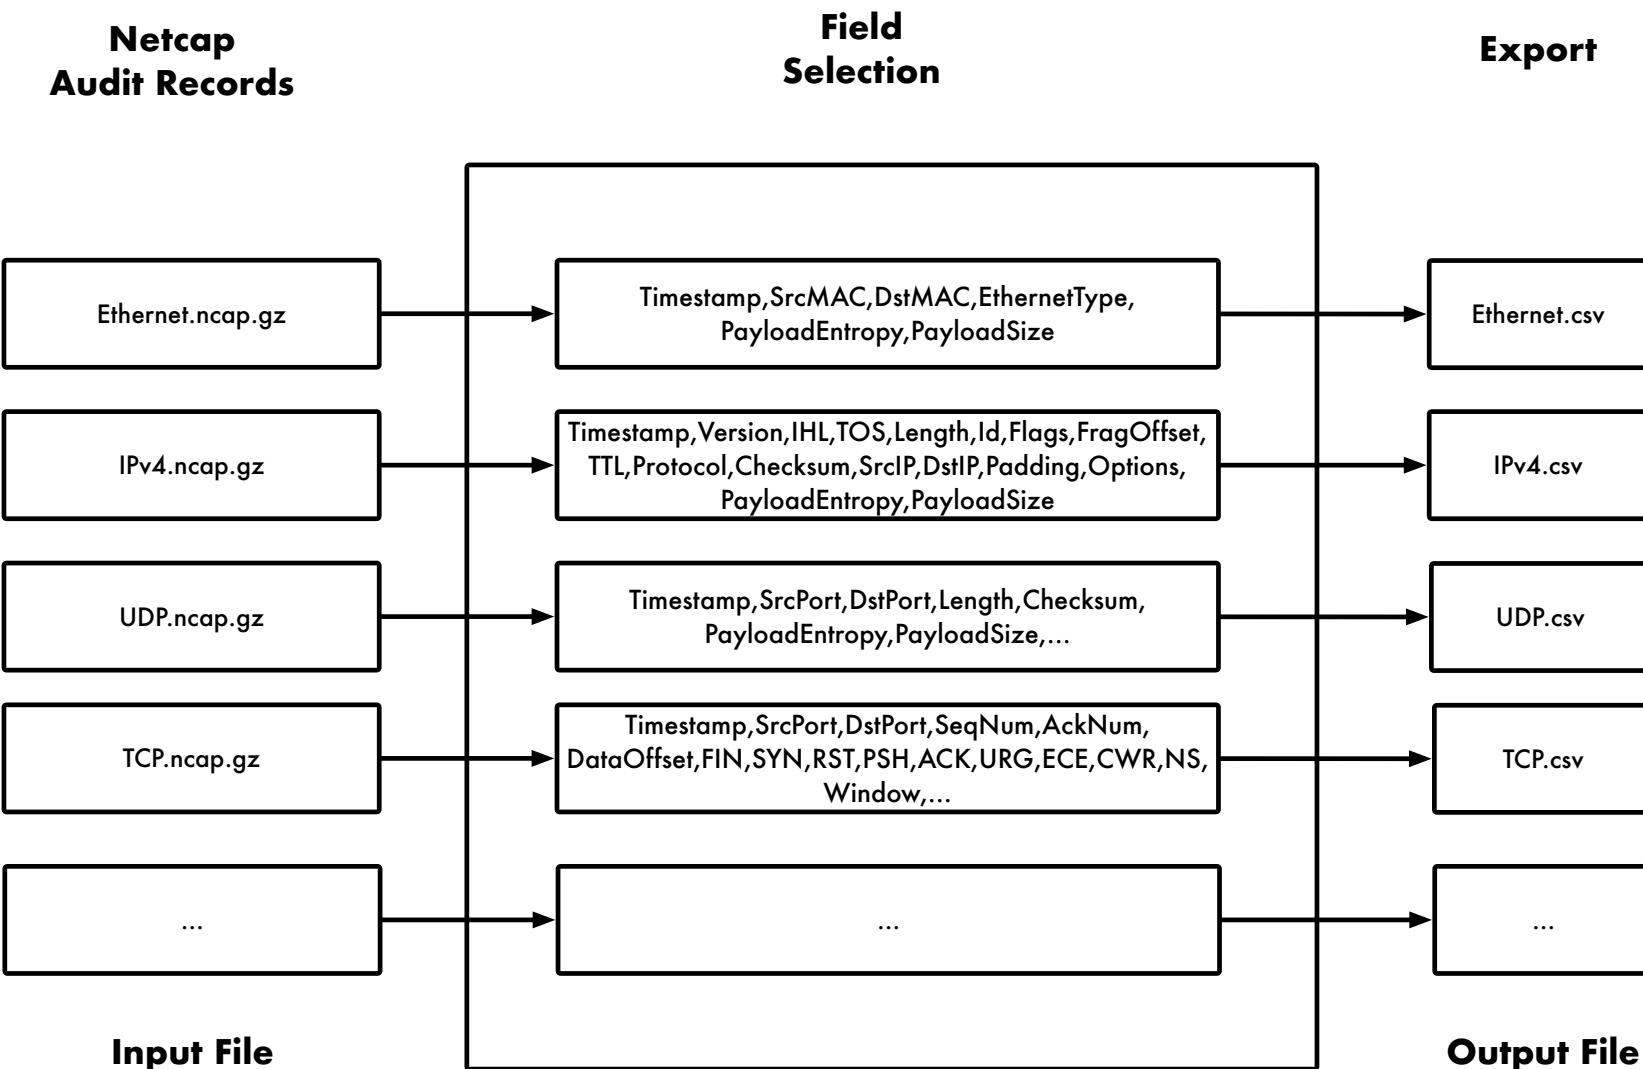
- + NTP
- + SIP
- + IGMP
- + LLC
- + IPv6HopByHop
- + SCTP
- + SNAP
- + LinkLayerDiscovery
- + ICMPv6NeighborAdvertisement
- + ICMPv6RouterAdvertisement
- + EthernetCTP
- + EthernetCTPReply
- + LinkLayerDiscoveryInfo

- + TCP
- + UDP
- + IPv4
- + IPv6
- + DHCPv4
- + DHCPv6
- + ICMPv4
- + ICMPv6
- + ICMPv6Echo
- + ICMPv6NeighborSolicitation
- + ICMPv6RouterSolicitation
- + DNS
- + ARP
- + Ethernet
- + Dot1Q
- + Dot11

Tell me more. **NETCAP** in a nutshell

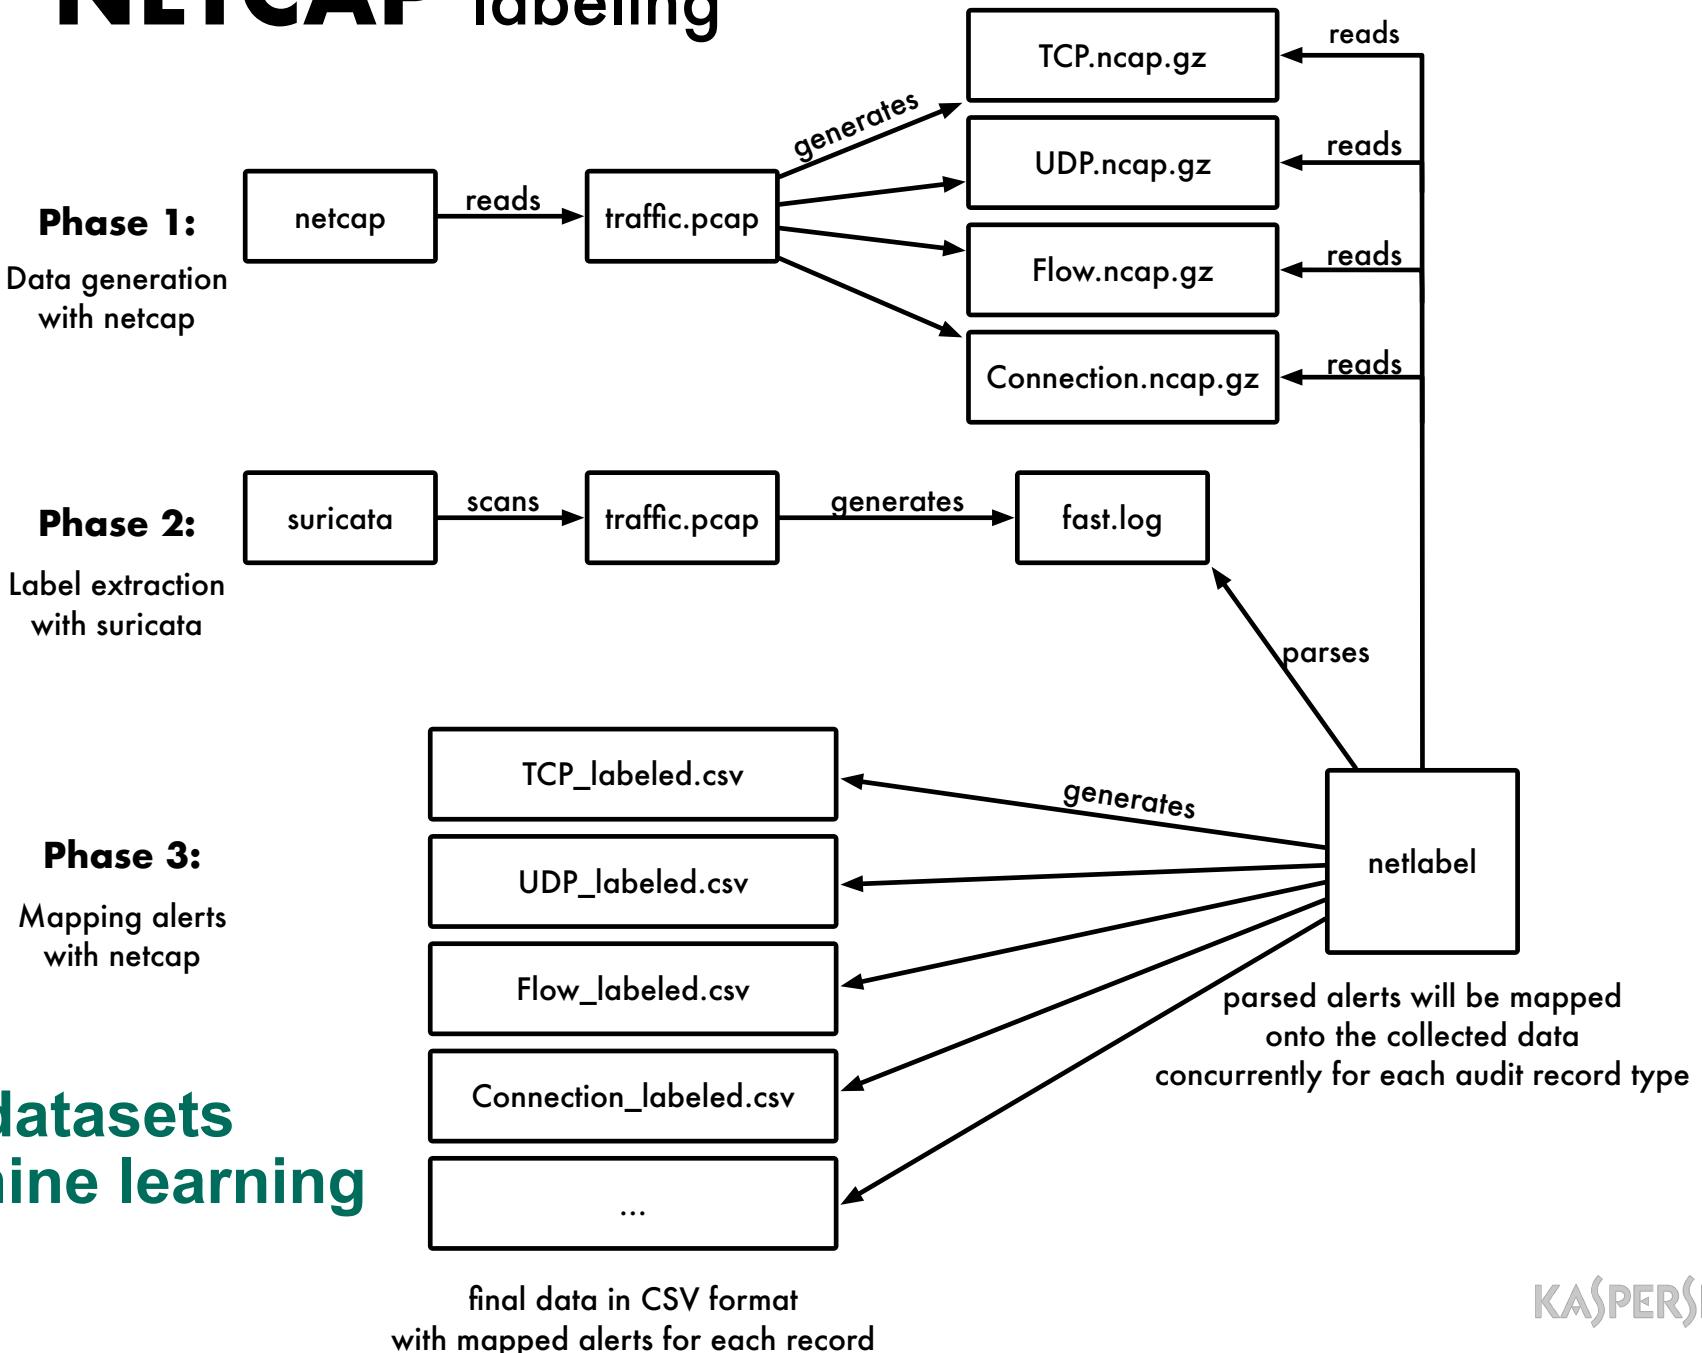


NETCAP filtering & csv export



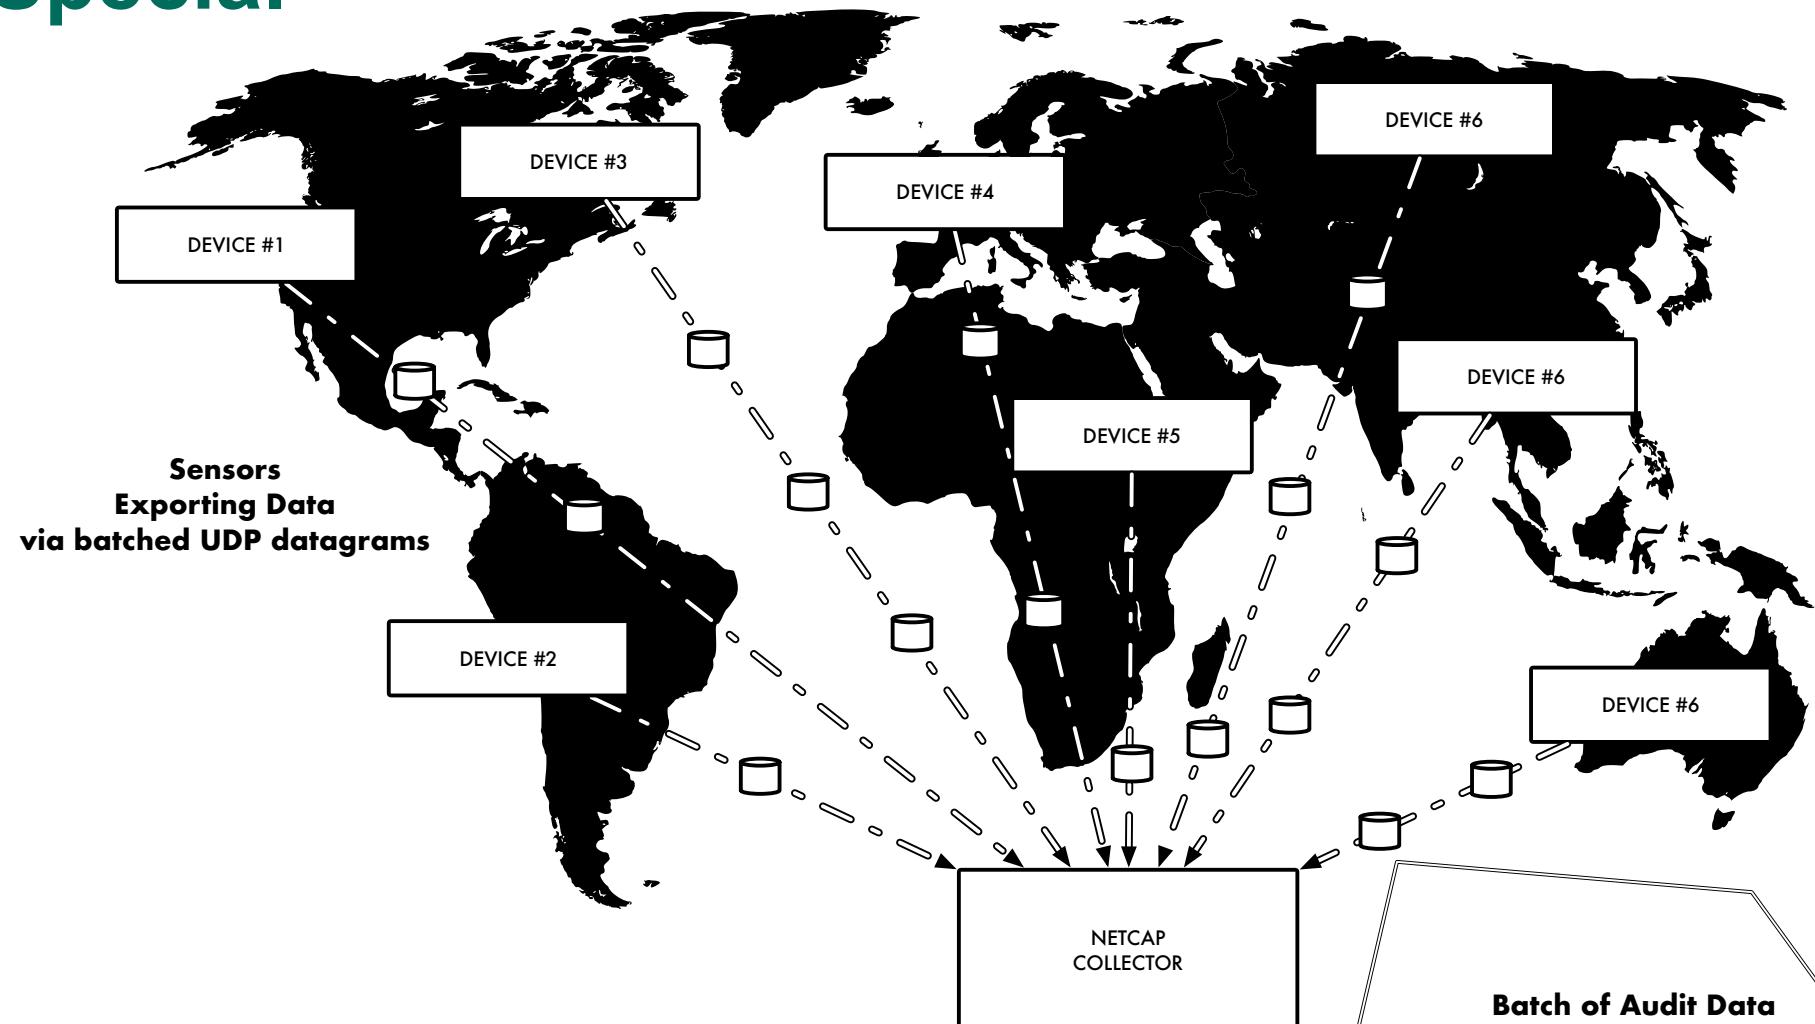
NETCAP labeling

Creation of labeled datasets for supervised machine learning



Cup Special

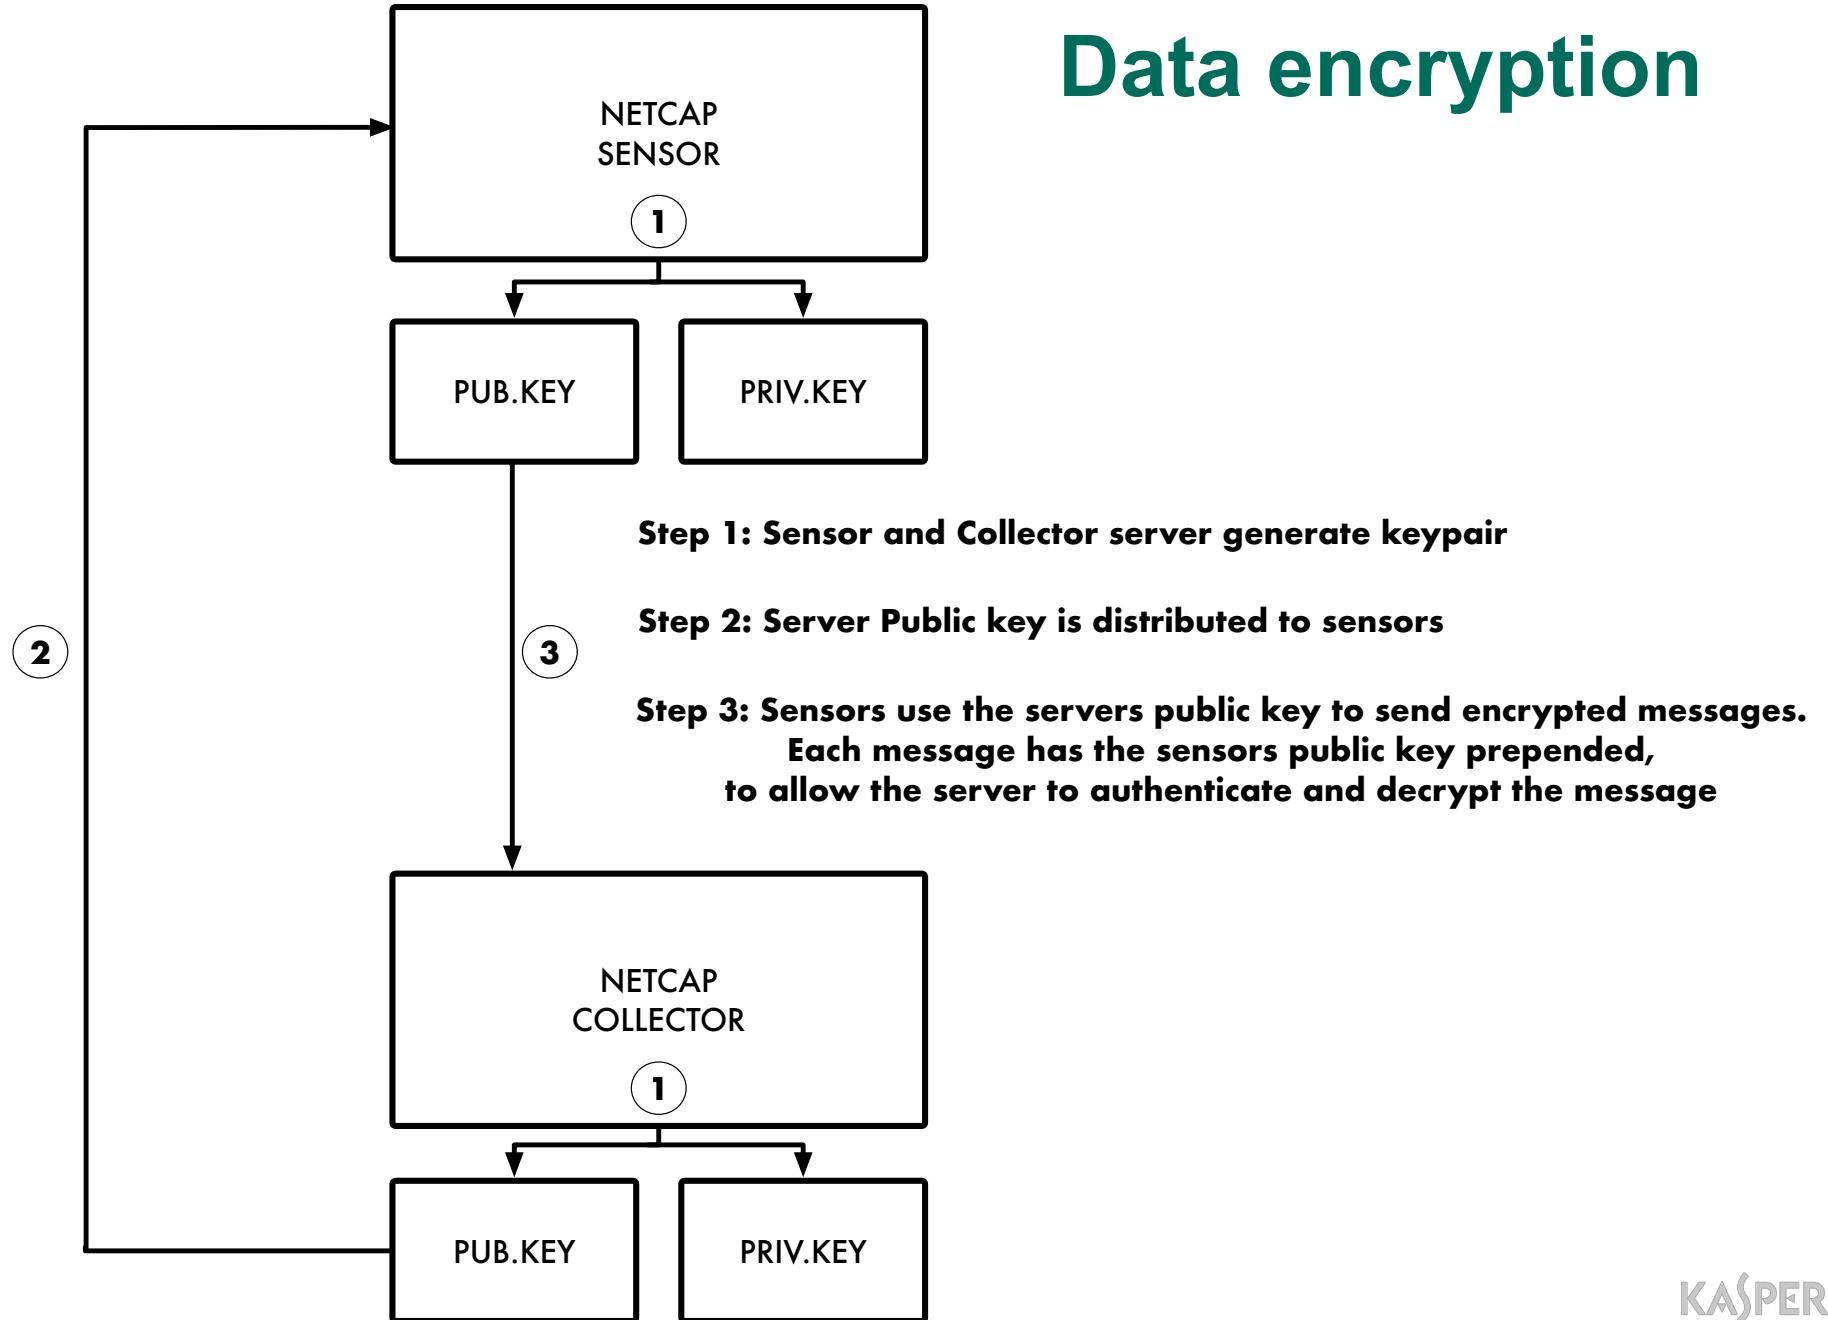
NETCAP Sensors



Distributed monitoring,
Traffic is encrypted in transit

Cup Special

Data encryption



Why monitor IoT devices?

A majority is compromised by brute force attacks - easy to detect with monitoring

Embedded applications are written in low level languages (C / C++) and may contain exploitable bugs

Many devices are infrequently updated and are vulnerable for a long period of time

State of IoT 2018?



From home and business security, industrial equipment and healthcare, to air travel, insurance and entertainment, the Internet of Things has spread rapidly across business sectors. Node.js is especially useful in IoT as this field continues to grow in popularity. Anytime a user has endpoints running Node.js in the field — whether it be in chipset or a kiosk/device — they need to ensure their apps are running, and doing so effectively. Collecting all that endpoint data effectively, and being able to repurpose it as well - is critical for the success of IoT implementations.

Traditionally, it's been hard to debug hardware. Writing software that enables hardware to "phone home" to web-based services can be especially tricky. Node.js is being adopted in IoT use cases for rapidly prototyping before products get hardened and mass produced. However, troubleshooting open source Node.js by itself in a debugging environment or as part of quality assurance can be difficult, if not impossible.

Wait... Node.js on IoT devices?



Sounds like a bad idea.



Casper Beyer [Follow](#)

A slightly bitter bearded developer with a background in arts.
Mar 30 · 4 min read

[Home](#) > [Security](#)

NEWS ANALYSIS

Malicious code in the Node.js npm registry shakes open source trust model

Bad actors using typo-squatting place 39 malicious packages in npm that went undetected for two weeks. How should the open source community respond?



By Fahmida Y. Rashid

Contributor, CSO | AUG 8, 2017 4:03 AM PT

Lucky for me, we live in an age where people install npm packages like they're popping pain killers.



David Gilbertson [Follow](#)

I like web stuff and cryptocurrency stuff.
Jan 6 · 10 min read

I'm harvesting credit card numbers and passwords from your site. Here's how.



Abdulrasaq Nasirudeen [FOLLOW](#)

Full stack software developer, Security researcher and an excellent communicator. I am highly solution o...

Few ways I could hijack your node.js applications

Published Feb 26, 2018 Last updated Aug 25, 2018

23



OpSecX

[MORE BY OPSECX](#)

The Node.js Ecosystem Is Chaotic and Insecure

Software

How one developer just broke Node, Babel and thousands of projects in 11 lines of JavaScript

Code pulled from NPM – which everyone was using

By Chris Williams, Editor in Chief 23 Mar 2016 at 01:24

167 SHARE ▾

Exploiting Node.js deserialization bug for Remote Code Execution

FEBRUARY 8, 2017 BLOG

tl;dr

Untrusted data passed into `unserialize()` function in node-serialize module can be exploited to achieve arbitrary code execution by passing a serialized JavaScript Object with an Immediately Invoked Function Expression (IIFE).



The good news: Go is gaining popularity for IoT applications

The screenshot shows the homepage of the Gobot website. At the top, there's a dark navigation bar with the word "GOBOT" in orange on the left and links for "Docs", "Platforms", "Resources", "News", "Blog", and "Github" on the right. Below the navigation is a large, colorful illustration of a brown robot with a red antenna and a grey head, standing on a green grassy hill under a blue sky with white clouds. The robot has a small tree growing from its back. To the left of the robot, the text "Go, Robot, Go! Golang Powered Robotics" is displayed in large, bold, black font. Below this, a subtitle reads "Next generation robotics/IoT framework with support for 35 different platforms". A yellow button with the text "Start Now" in black is positioned on the left side of the illustration. At the bottom of the main section, there are social media icons for GitHub (4,879 stars), Fork (608 forks), Twitter (Tweet), and Twitter (Follow @gobatio). A dark, semi-transparent overlay at the bottom contains the text "Gobot is a framework for robots, drones, and the Internet of Things (IoT), written in the Go programming language" in white.

What about embedded systems?

Go depends on features of the operating system.

However, there is work in progress on bringing it to embedded devices as well.

Also the go authors are aware of this - maybe Go 2 will change the game?

README.md

TinyGo - Go compiler for microcontrollers

docs passing build passing

We never expected Go to be an embedded language and so it's got serious problems [...].

-- Rob Pike, [GopherCon 2014 Opening Keynote](#)

TinyGo is a project to bring Go to microcontrollers and small systems with a single processor core. It is similar to [emgo](#) but a major difference is that I want to keep the Go memory model (which implies garbage collection of some sort). Another difference is that TinyGo uses LLVM internally instead of emitting C, which hopefully leads to smaller and more efficient code and certainly leads to more flexibility.

My original reasoning was: if [Python](#) can run on microcontrollers, then certainly [Go](#) should be able to and run on even lower level micros.

Whats your business model?

GPLv3 :)

Further open source contributions

google / gopacket

Watch 125 Unstar 2,299 Fork 437

Code Issues 56 Pull requests 13 Projects 0 Insights

reasemblydump example: Fixed data races on errorsMap and tcpStream #560

Merged gconnell merged 1 commit into google:master from dreadl0ck:master 14 days ago

Conversation 1 Commits 1 Checks 0 Files changed 1 +12 -1

dreadl0ck commented 14 days ago

The reasemblydump example contains two data races:
The errorsMap is potentially accessed simultaneously by multiple goroutines.
A mutex was added to make it thread safe.

When using stream.run for both client and server for parsing http requests and responses,
both use a pointer to the parent stream to access the stream.urls array without any synchronization.
A mutex was also added to the tcpStream structure to prevent this.

reasemblydump example: Fixed data races on errorsMap and tcpStream 1463509

gconnell merged commit ec90f6c into google:master 14 days ago

View details Revert

2 checks passed

Reviewers
No reviews

Assignees
No one assigned

Labels
None yet

Projects
None yet

Milestone
None yet

KASPERSKY

Further open source contributions

 [dreadlOck / ja3](#) Watch ▾ 2 Star 30 Fork 2

[Code](#) [Issues 0](#) [Pull requests 0](#) [Projects 0](#) [Wiki](#) [Insights](#) [Settings](#)

Ja3 TLS Client Hello Hashes in Go Edit

 [dreadlOck / gopcap](#) Watch ▾ 0 Star 28 Fork 0

[Code](#) [Issues 0](#) [Pull requests 0](#) [Projects 0](#) [Wiki](#) [Insights](#) [Settings](#)

Fast Golang PCAP Reader & Benchmark Comparison Edit

 [dreadlOck / cryptoutils](#) Watch ▾ 0 Star 0 Fork 0

[Code](#) [Issues 0](#) [Pull requests 0](#) [Projects 0](#) [Wiki](#) [Insights](#) [Settings](#)

A thin wrapper around the NaCl toolkit and a few utility functions Edit

Other use cases?

Monitor honeypots

Monitor medical devices

Forensic Analysis

Research! :)

Core benefits of netcap?

Memory Safety (pure Golang codebase)

Data Accessibility (Protocol Buffers)

Performance (all 40+ encoders, ~1GB/min processing time)

Documentation (2k lines of comments on 8k lines of Go code)

130+ pages thesis with in depth technical documentation and usage

The screenshot shows the Canadian Institute for Cybersecurity website. At the top, there's a black header with the UNB logo and navigation links for 'Give to UNB', 'Apply', and a search icon. Below the header is a red navigation bar with links for 'Home', 'About', 'Research', 'Members', and 'Datasets'. On the left side, there's a sidebar titled 'Datasets' with a list of datasets: 'IDS 2012', 'IDS 2017' (which is highlighted with a red border), 'NSL-KDD', 'VPN-nonVPN', 'Botnet', 'Android Validation', 'Android Botnet', 'Tor-nonTor', 'Dos Dataset', 'Android-Adware', 'Android-Malware2017', and 'CSE-CIC-IDS2018'. The main content area features a large title 'Intrusion Detection Evaluation Dataset (CICIDS2017)' and three paragraphs of descriptive text.

Intrusion Detection Evaluation Dataset (CICIDS2017)

Intrusion Detection Systems (IDSs) and Intrusion Prevention Systems (IPs) are the most important defense tools against the sophisticated and ever-growing network attacks. Due to the lack of reliable test and validation datasets, anomaly-based intrusion detection approaches are suffering from consistent and accurate performance evolutions.

Our evaluations of the existing eleven datasets since 1998 show that most are out of date and unreliable. Some of these datasets suffer from the lack of traffic diversity and volumes, some do not cover the variety of known attacks, while others anonymize packet payload data, which cannot reflect the current trends. Some are also lacking feature set and metadata.

CICIDS2017 dataset contains benign and the most up-to-date common attacks, which resembles the true real-world data (PCAPs). It also includes the results of the network traffic analysis using CICFlowMeter with labeled flows based on the time stamp, source and destination IPs, source and destination ports, protocols and attack (CSV files).

Generating realistic background traffic was our top priority in building this dataset. We have used our proposed B-Profile system (Sharafaldin, et al. 2016) to profile the abstract behaviour of human interactions and generates a naturalistic benign background traffic. For this dataset we built the abstract behaviour of 25 users based on the HTTP, HTTPS, FTP, SSH, and email protocols.

**Up-to-date, well documented dataset
~50GB pcap dumps**

Tuesday-WorkingHours.pcap

File	Num Records	Size	Labels	Exec Time	Validation Score
Connection_labeled.csv	284166	51M	1809	1m 14s	0.9936234903296641
DNS_labeled.csv	700447	144M	37	2m 3s	0.999965736214537
Ethernet_labeled.csv	11551954	880M	3551	32m 24s	0.999699098576899
Flow_labeled.csv	647294	112M	3340	1m 57s	0.9946855843385406
HTTP_labeled.csv	45800	14M	4214	39s	0.9363318777292576
NTP_labeled.csv	15507	3.0M	13	5s	0.9987103430487491
NetworkFlow_labeled.csv	29094	4.0M	1292	9s	0.9572449821281276
TCP_labeled.csv	10710230	1.5G	3450	30m 59s	0.9996720892694014
TransportFlow_labeled.csv	212613	20M	1721	43s	0.99153403318659
UDP_labeled.csv	787015	43M	46	2m 9s	0.9999390101344826

Table 6.35: Classification results experiment 6 Tuesday-WorkingHours.pcap

Thesis? Coming soon!

INSTITUT FÜR INFORMATIK

DER LUDWIG-MAXIMILIANS-UNIVERSITÄT MÜNCHEN



Bachelorarbeit

Concept and implementation of memory safe feature collection for network intrusion detection

Conclusion

Use memory safe programming languages for the development of critical software infrastructure

This allows us to focus on the creation of solid logic, rather than solid memory management.

Golang || Rust



SECUR'IT CUP'18

Thank you!
Questions?

philipp.mieden@protonmail.ch
github.com/dreadl0ck
twitter.com/dreadcode